

Object Detection

Slow Fast Faster RCNN

Vision@OUC

Wang Chao

Group of DL

Overview

- Introduction
- Slow Fast Faster R-CNN
- Py-faster-rcnn
- Other works
- Q&A

Introduction

To give computers visual intelligence

---Feifei Li



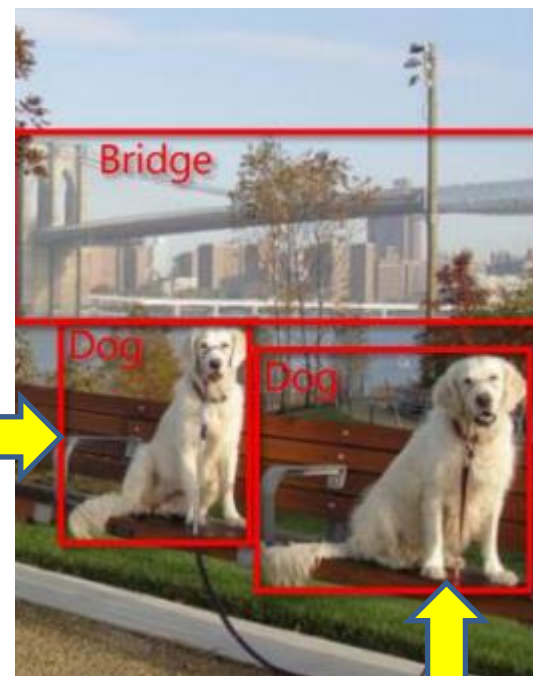
Classification VS Detection

Classification: What

Detection: What and Where



Localization
Where ?



Recognition
What ?



Usually, We need a **bounding box** to tell us “What” and “Where”

Efficient Object Detection

- Object detection is arguably a harder problem than image classification
- Usually a large number of image sub-window need to scanned in order to **localize objects**, leading to heavy computational processing
- Challenge: In many real-world applications, running a **fast** object detector is as critical as running an **accurate** object detector

What is an elegant object detector

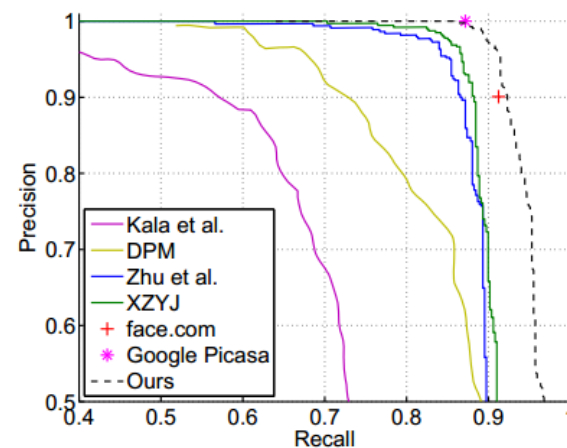
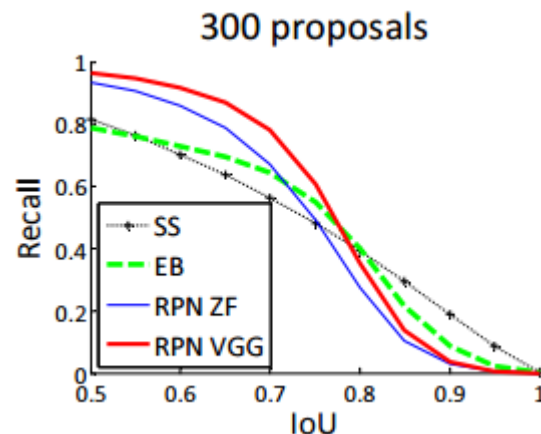
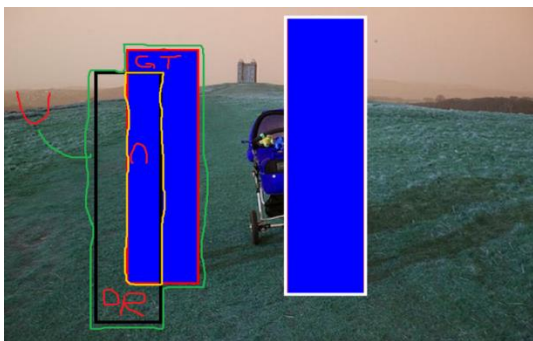
- Speed:
 - Per image
- Accuracy
 - MAP
- Others
 - Training time
 - Model
 - generalization

Evaluation

- Box match
 - IoU: Intersection-over-Union

$$IOU = \frac{DetectionResult \cap GroundTruth}{DetectionResult \cup GroundTruth}$$

- Accuracy
 - Recall
 - Miss rate
 - False Positive per Image
 - Precision
 - Average Precision



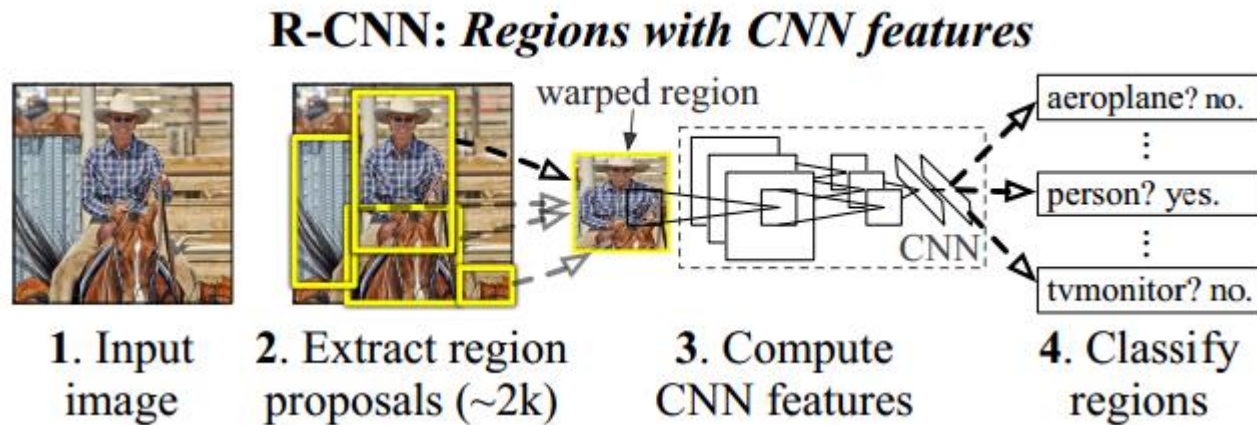
- Dataset
 - Pascal VOC
 - 2007, 2010, 2012
 - 20 categories
 - Ms COCO
 - 2015
 - 80 categories
 - ImageNet
 - ILSVRC 2013,2014,2015,2016
 - 200 categories
 - Others
 - Face: AFW, FDDB, MAF, IJB-A
 - Pedestrian : INRIA, KITTI
 - Vehicle: KITTI

R-CNN

- Past methods: complex **ensemble systems**
 - Combine multiple low-level image features with high-level context
- RCNN: Regions with CNN features
 - CNNs for region proposal
 - Supervised pre-training for **scarce** training data

Features matter !

- R-CNN architecture



- Takes an input image
- Extracts around **2k** bottom-up region proposals
- Compute features for **each proposal** using CNN
- Classifies each proposal region using linear SVMs

Training pipeline

Steps for training a slow R-CNN detector

1. [offline] $M \lll$ **Pre-train** a ConvNet for ImageNet Classification
2. $M' \lll$ **Fine-tune** M for object detection (softmax + log loss)
3. $F \lll$ Cache feature vectors to disk using M'
4. Train **post hoc linear SVMs** on F (hinge loss/L2 loss)
5. Train post hoc linear bounding-box regressors on F (squared loss)

* Ignoring pre-training, there are three separate training stages

Why I call it “Slow” R-CNN

Example timing for R-CNN on VOC07(only 5k training images) using VGG16 and a K40 GPU

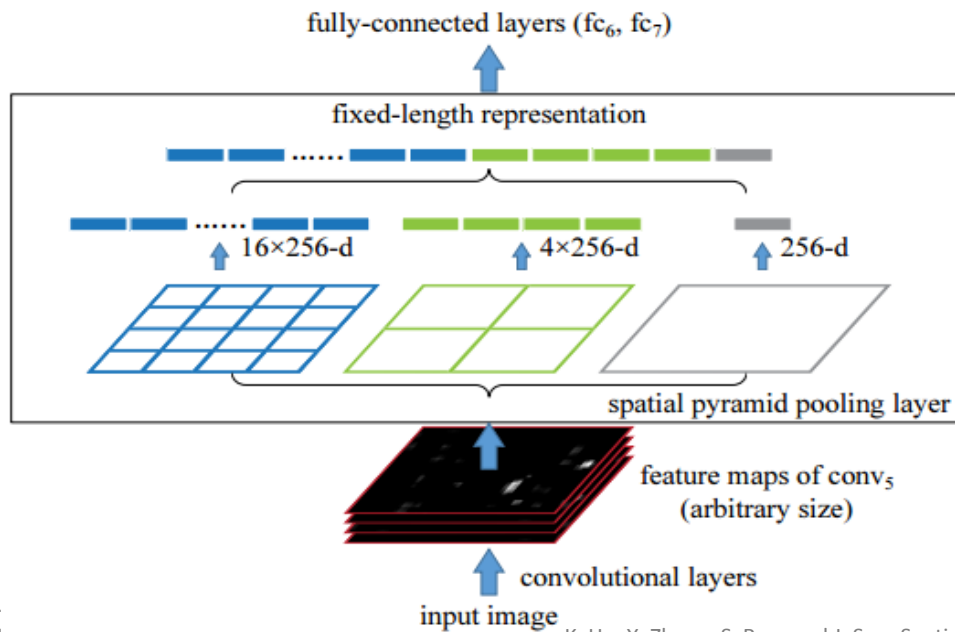
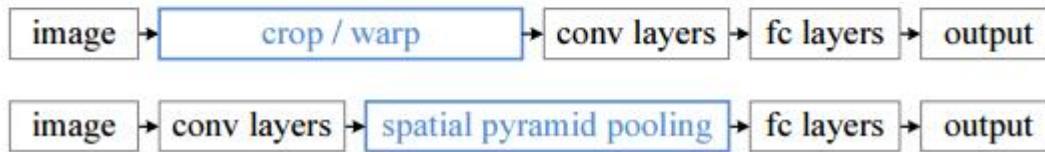
Fine-tune(BP, SGD): 18 hours

Feature extraction: 63 hours

SVM and bounding-box regressor training: 3 hours

Total: 84 hours

SPP-net



SPP-net helps

Example timing for R-CNN/SPP-net

- Fine-tuning(BP, SGD): 18 hours/16 hours
- Feature extraction: 63 hours/ 5.5 hours
 - Forward pass time(helps here)
 - Disk I/O is costly(dominates SPP-net extraction time)
- SVM and bounding-box regressor training 3 hours/4 hours
- Total: 84 hours/25.5 hours

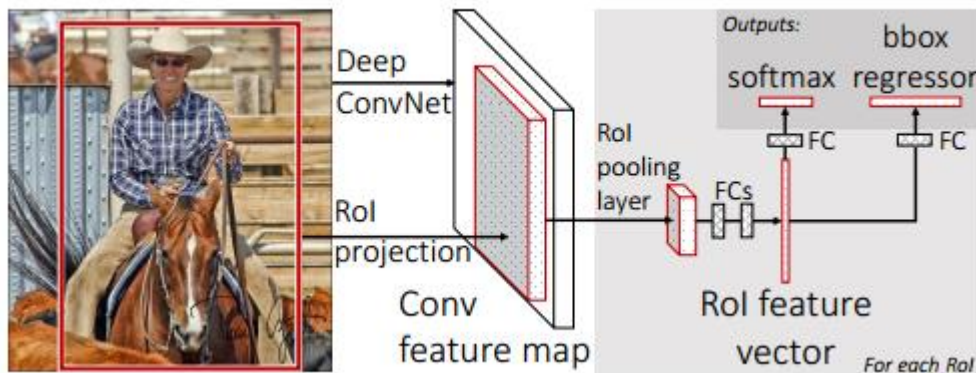
Fast R-CNN

Train models in multi-stage pipelines that are **slow** and **inelegant**

R-CNN and SPPnet

1. Training is a **multi-stage** pipeline
ConvNet; SVM; bounding-box regressor
2. Training is expensive in **space and time**
features of each propos in each image
3. Object detection is **slow**
features extract ; VGG16 47s/image(GPU)

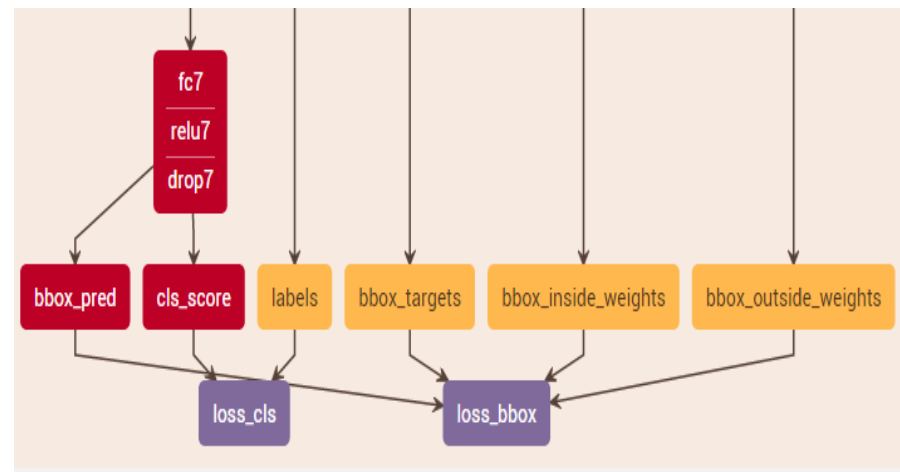
Fast R-CNN architecture



- An input image and multiple **Rols** are input into fully convNet
- Each Rols is pooled into a **fixed-size feature map** and then mapped to a feature vector by FC layes
- The network has two output per RoI: **Softmax probability and bouding-box regression**
- Trained end-to-end with **multi-task loss**

Training Fast R-CNN end-to-end

- Define **one network** with two loss branches
 - Branch 1: softmax classifier
 - Branch 2: linear bounding-box regressors
 - Overall loss is the sum of the two loss branches
- Fine-tune the network jointly with SGD
 - Optimizes features for both tasks
- BP errors all the way back to the conv layers



Benefits of end-to-end training

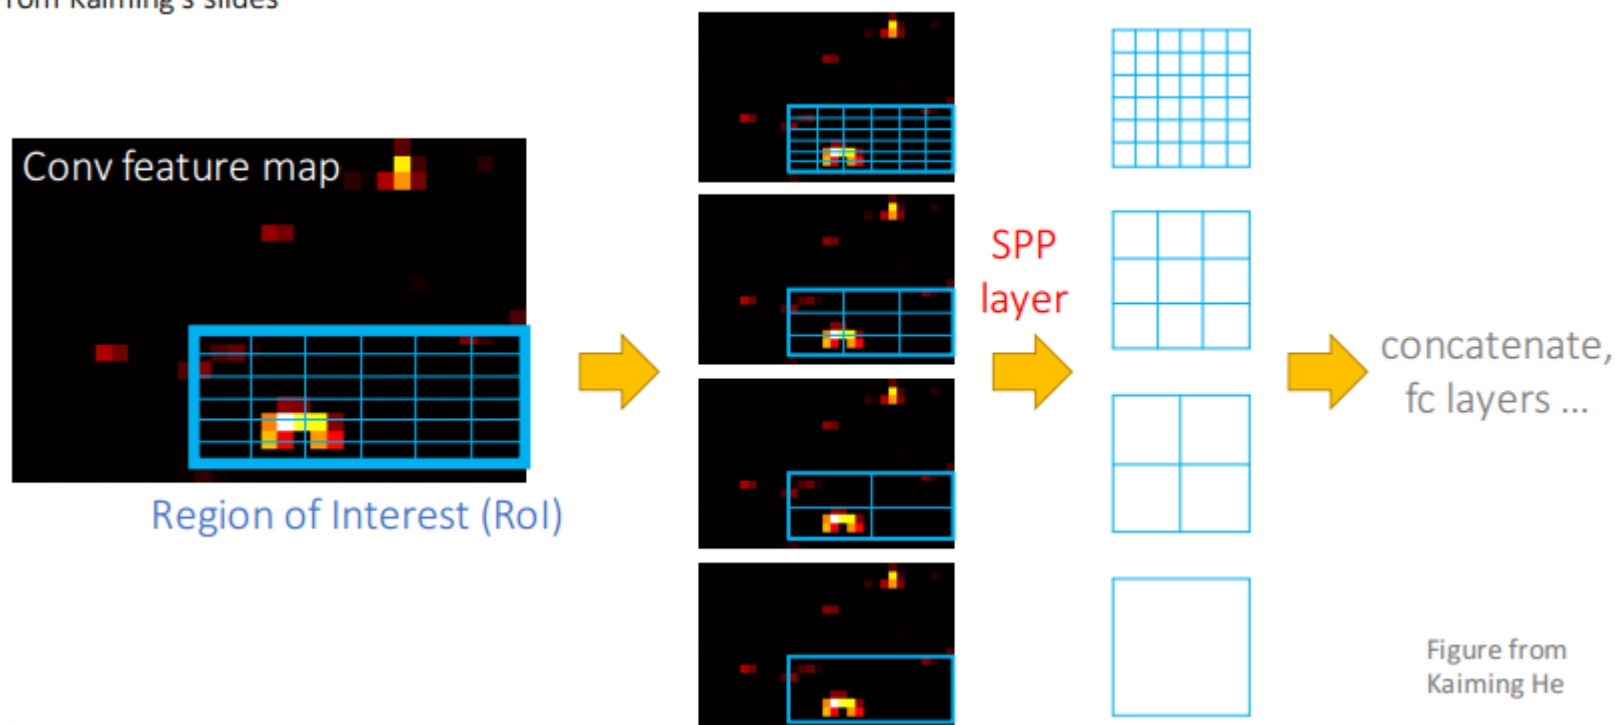
- Simpler implementation
- Faster training
 - No reading/writing features from/to disk
 - No training post hoc SVMs and bounding-box regressors
- Optimizing a **single multi-task objective** work better than optimizing objectives independently

End to End training requires overcoming two technical obstacles

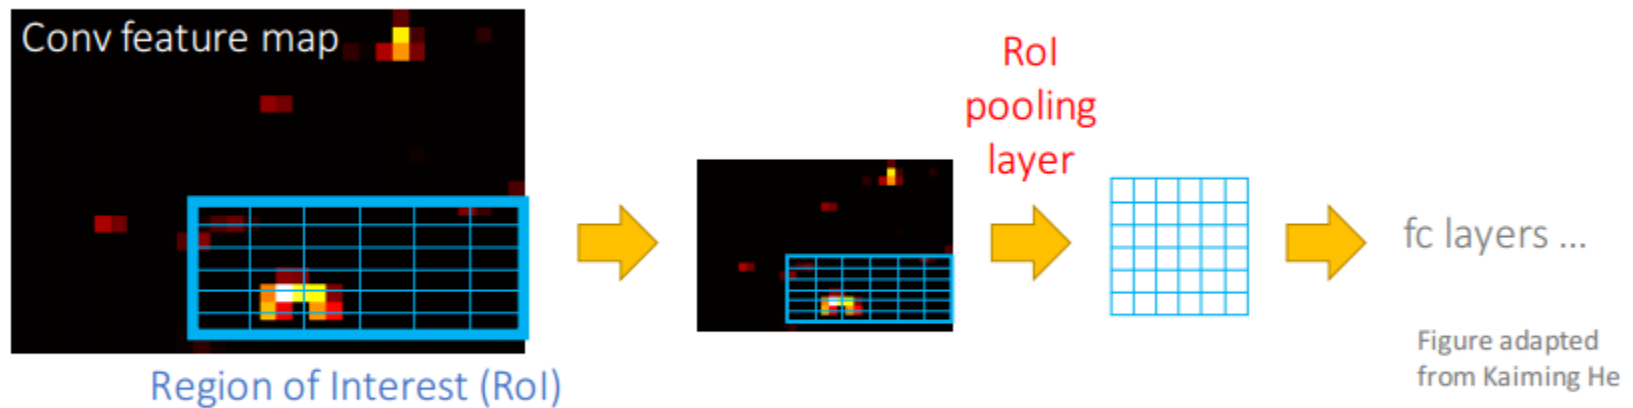
Obstacle1: ROI pooling

- Spatial Pyramid Pooling layer

From Kaiming's slides



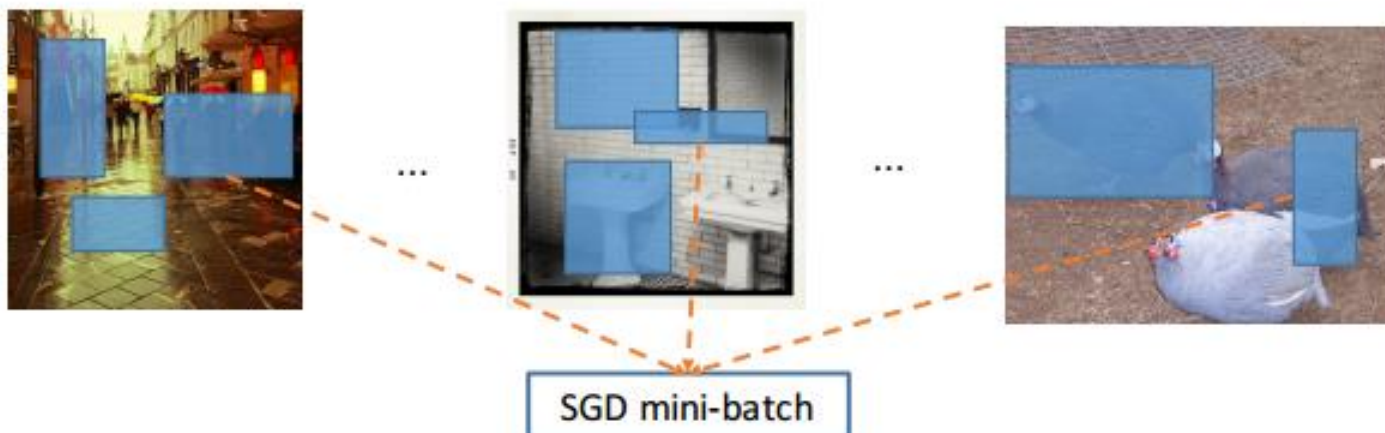
Region of Interest pooling layer



Just a special case of the SPP layer with one pyramid level

Obstacle2: Make SGD steps efficient

- R-CNN and SPP-net use region-wise sampling to make mini-batch
 - Sample 128 example ROIs uniformly at random
 - Examples will come from different images



- **Solution:** use hierarchical sampling to build mini-batches
 - Sample a **small number** of images(2)
 - Sample **many examples** from each image



Cost per mini-batch compared to slow R-CNN

$$\frac{\text{input size for Fast R-CNN}}{2 * 600 * 1000} / \frac{\text{input size for slow R-CNN}}{(128 * 224 * 224)} = 0.19x < \text{computation than slow R-CNN}$$

Towards fast and end-to-end

- Multi-task loss for unified network

$$L(p, u, t^u, v) = L_{\text{cls}}(p, u) + \lambda[u \geq 1]L_{\text{loc}}(t^u, v),$$

- Truncated SVD for faster detection

Compress FC layers

$$W \approx U\Sigma_t V^T$$

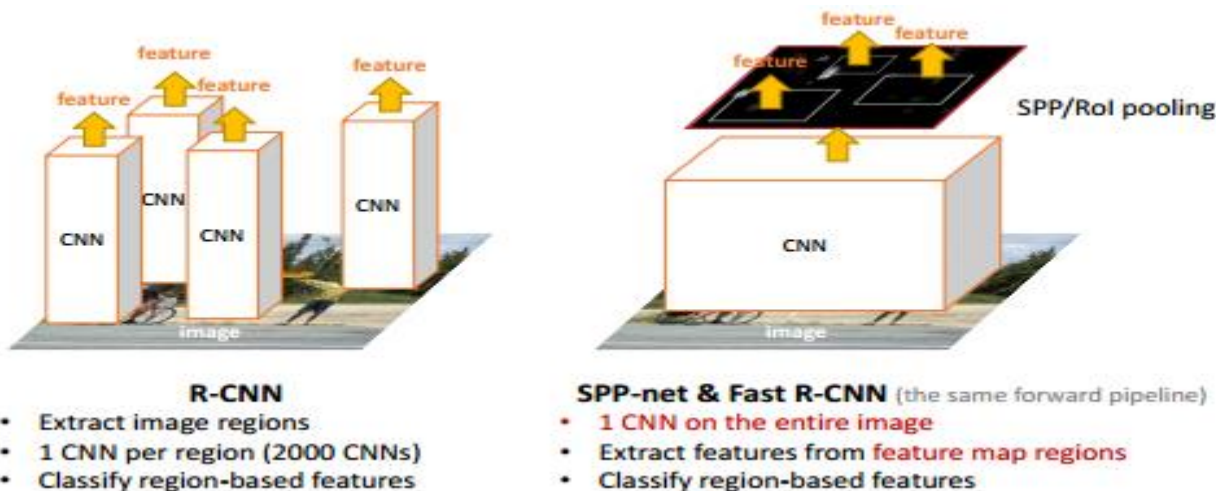
Before: $u \times v$

After: $t(u + v)$

Especially: $t \leq \min(u, v)$

Fast R-CNN outcome

- Better training time and testing time with better accuracy than slow R-CNN / SPP-net
- Training time: 84 hours/25.5hours/8.75hours
- VOC07 test mAP: 66.0%/63.1%/68.1%
- Testing time per image: 47s/2.3s/0.32s
 - Plus 0.2 to > 2s per image depending on proposal method
 - With selective search: 49s/4.3s/2.32s



Faster R-CNN

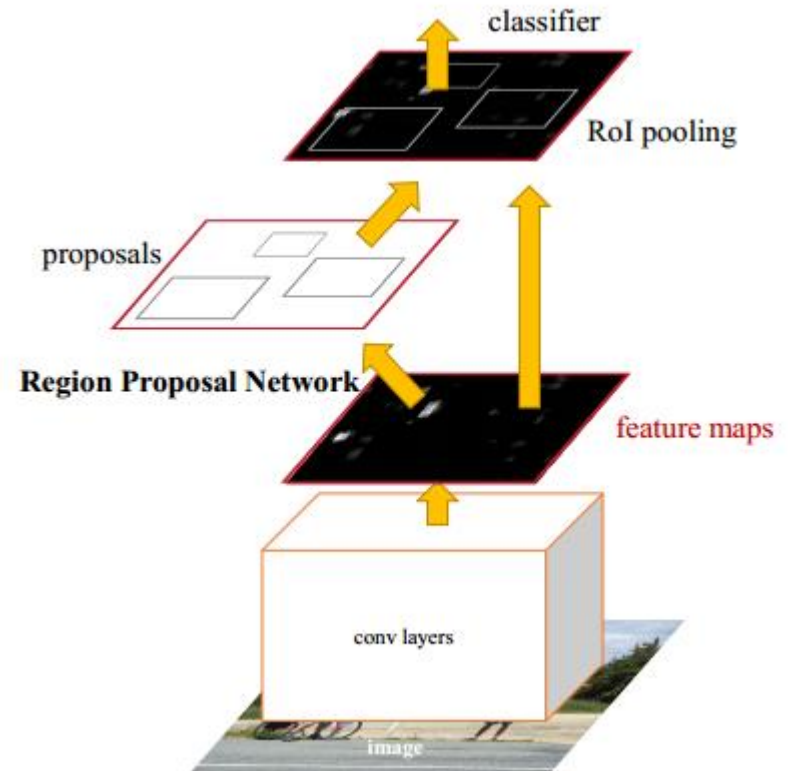
State-of-the-art object detection networks depend on **region proposal algorithms** to hypothesize object locations. Advances like SPP-net and Fast R-CNN have reduced the running time, **exposing region proposal computation as a bottleneck**.

Faster R-CNN architecture

Faster R-CNN =
RPN + Fast RCNN

Does not depend on an
external region proposal
algorithm

Does object detection in a
single forward pass



Training method

1. Train RPN: RPN is initialized with imagenet pre-trained model and fine-tuned end to end for region proposal task
2. Use region proposal by RPN, train Fast-RCNN : initialized by ImageNet pre-trained model
3. Train RPN and fix the shared convolutional layers
4. Train Fast R-CNN and fix the convolutional layers and form a unified network.

Region Proposal Network

To generate region proposals:

Slide window in last shared ConvNet: 3×3

Lower-dimensional feature: 256D(FZ)\512D(VGG16)

Generate: cls loss and reg loss

1×1 convolutional layer

Anchors:

Multibox uses k-means: 800 anchors

Multibox: $(4 + 1) \times 800$ dimensional fc output layer

Output layer:

6.1×10^6 , $(1536 \times (4 + 1) \times 800)$ for googlenet

RPN: $(4 + 2) \times 9$ dimensional conv output layer

output layer: 2.8×10^4 ,

$512 \times (4 + 2) \times 9$ for VGG16

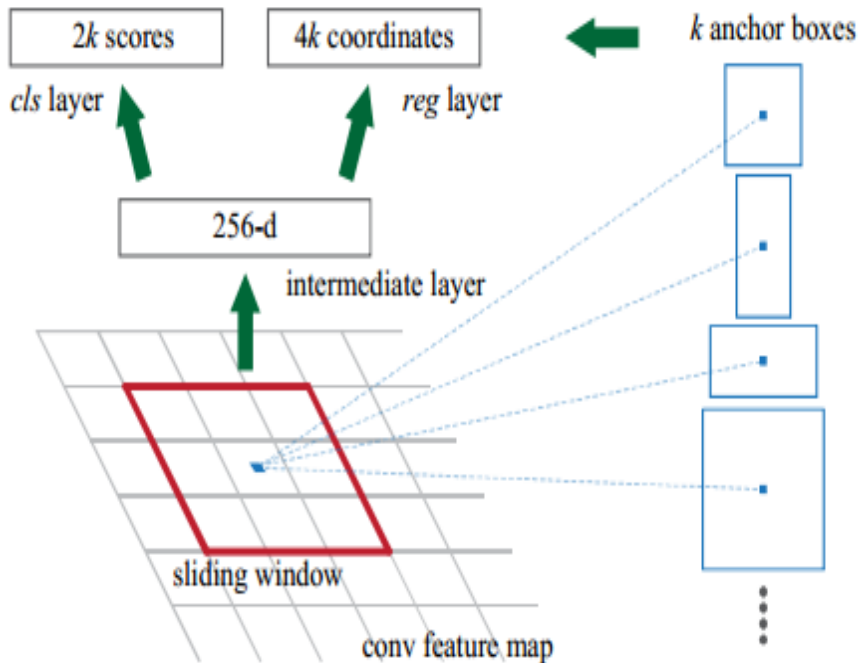
Multi-scale design:

From single-scale image; 3 scale and 3 aspect ratios

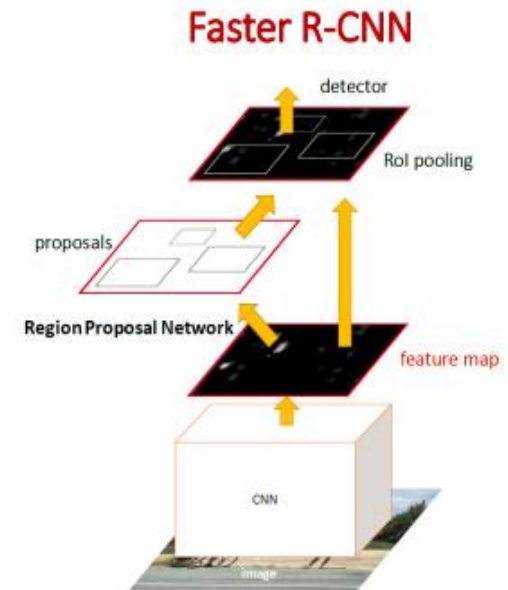
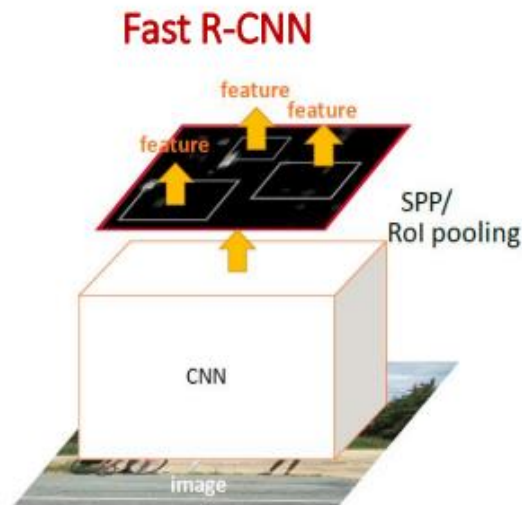
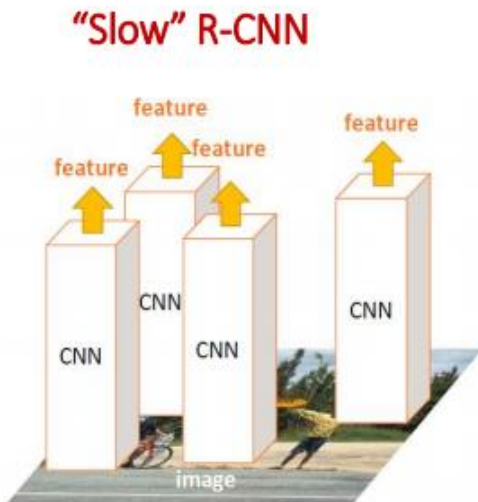
Scale: 128, 256, 512; ratios: 1:1, 1:2, 2:1

Translation-Invariant Anchors

Without extra cost for addressing scales



Slow vs Fast vs Faster



Py-Faster-RCNN 环境搭建

Requirements: hardware

Installation

1. Clone the Faster R-CNN

```
git clone --recursive https://github.com/rbgirshick/py-faster-rcnn.git
```

2. Build Cython modules

```
cd $FRCN_ROOT/lib
```

```
Make
```

3. Build Caffe and pycaffe(see: Caffe installation instructions)

Note: In Makefile.config

```
WITH_PYTHON_LAYER := 1
```

4. Download pre-computed Faster R-CNN detectors

```
cd $FRCN_ROOT
```

```
./data/scripts/fetch_faster_rcnn_models.sh
```

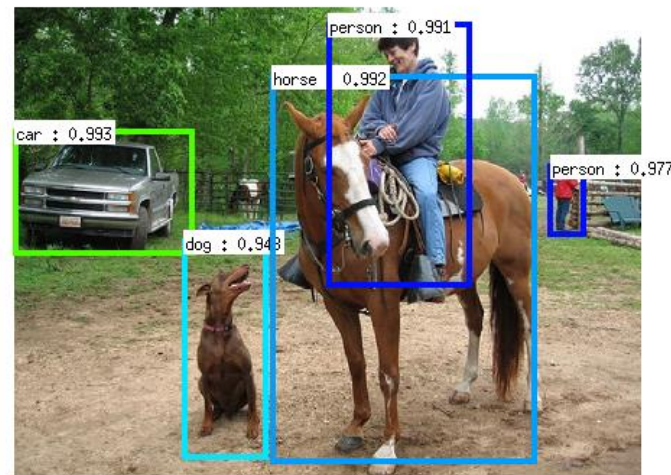
5. Demo

After successfully completing basic installation

To run the demo

```
cd $FRCN_ROOT
```

```
./tools/demo.py
```



Questions occurred

- cudnn version problems

For **cudnn V5**:

1. Replace these document with new caffe

`include/caffe/layers/cudnn_relu_layer.hpp,`

`src/caffe/layers/cudnn_relu_layer.cpp, src/caffe/layers/cudnn_relu_layer.cu`

`include/caffe/layers/cudnn_sigmoid_layer.hpp,`

`src/caffe/layers/cudnn_sigmoid_layer.cpp, src/caffe/layers/cudnn_sigmoid_layer.cu`

`include/caffe/layers/cudnn_tanh_layer.hpp,`

`src/caffe/layers/cudnn_tanh_layer.cpp, src/caffe/layers/cudnn_tanh_layer.cu`

2. Replace the name in `include/caffe/util/cudnn.hpp`

`cudnnConvolutionBackwardData_v3` 函数名替换为 `cudnnConvolutionBackwardData`

`cudnnConvolutionBackwardFilter_v3` 函数名替换为 `cudnnConvolutionBackwardFilter`

- Demo in CPU version

http://blog.sina.com.cn/s/blog_679f93560102wpyf.html

More than Demo

- Datasets: MS coco/ ImageNet
 - factory.py
 - new dataset class: `my_dataset.py(pascal_voc.py)`
 - `_load_image_set_index`, `image_path_from_index`
 - `_load_pascal_annotation`
- Sources:
 - <https://github.com/deboc/py-faster-rcnn/tree/master/help>
 - <https://github.com/rbgirshick/py-faster-rcnn/issues/243>

New network architecture

- ResNet-152
 - Download the prototxt:

<https://github.com/KaimingHe/deep-residual-networks>
 - train.prototxt
 - test.prototxt
- Train your model
 - class number: input-data layer
 - Add Fast R-CNN layer
 - RPN layer

```
layer {
  name: "rpn_conv/3x3"
  type: "Convolution"
  bottom: "res4f"
  top: "rpn/output"
  param { lr_mult: 1.0 }
  param { lr_mult: 2.0 }
  convolution_param {
    num_output: 512
    kernel_size: 3 pad: 1 stride: 1
    weight_filler { type: "gaussian" std: 0.01 }
    bias_filler { type: "constant" value: 0 }
  }
}
```

```
...
name: "rpn_relu/3x3"
...
name: "rpn_cls_score"
...
name: "rpn_bbox_pred"
...
name: "rpn_cls_score_reshape"
...
name: rpn-data
...
name: "rpn_loss_cls"
...
name: "rpn_loss_bbox"
...
```

Other works recently

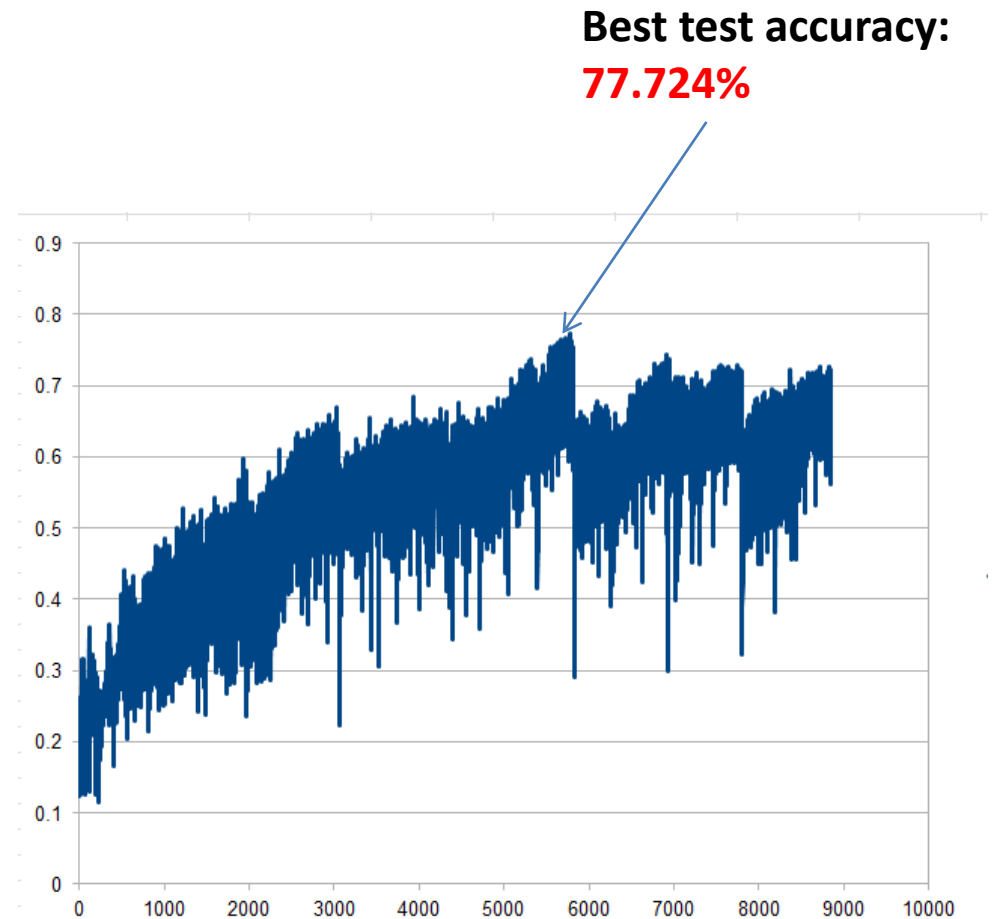
- LSTM **converge** & bidirectional LSTM
- Deploy sever with python Flask

LSTM

1. Error fixed

2. Momentum

3. L2 Norm (weight decay)



Flask for a web sever

A small application:

```
from flask import Flask
app = Flask(__name__)

@app.route('/')
def hello_world():
    return 'Hello World!'

if __name__ == '__main__':
    app.run()
```

```
import datetime
import logging
import flask
import werkzeug
import optparse
import tornado.wsgi
import tornado.httpserver
import numpy as np
import pandas as pd
from PIL import Image
import cStringIO as StringIO
import urllib
import exifutil
```

```
REPO_DIRNAME = os.path.abspath(os.path.dirname(os.path.abspath(__fi
UPLOAD_FOLDER = '/tmp/caffe_demos_uploads'
ALLOWED_IMAGE_EXTENSIONS = set(['png', 'bmp', 'jpg', 'jpe', 'jpeg',

# Obtain the flask app object
app = flask.Flask(__name__)
```

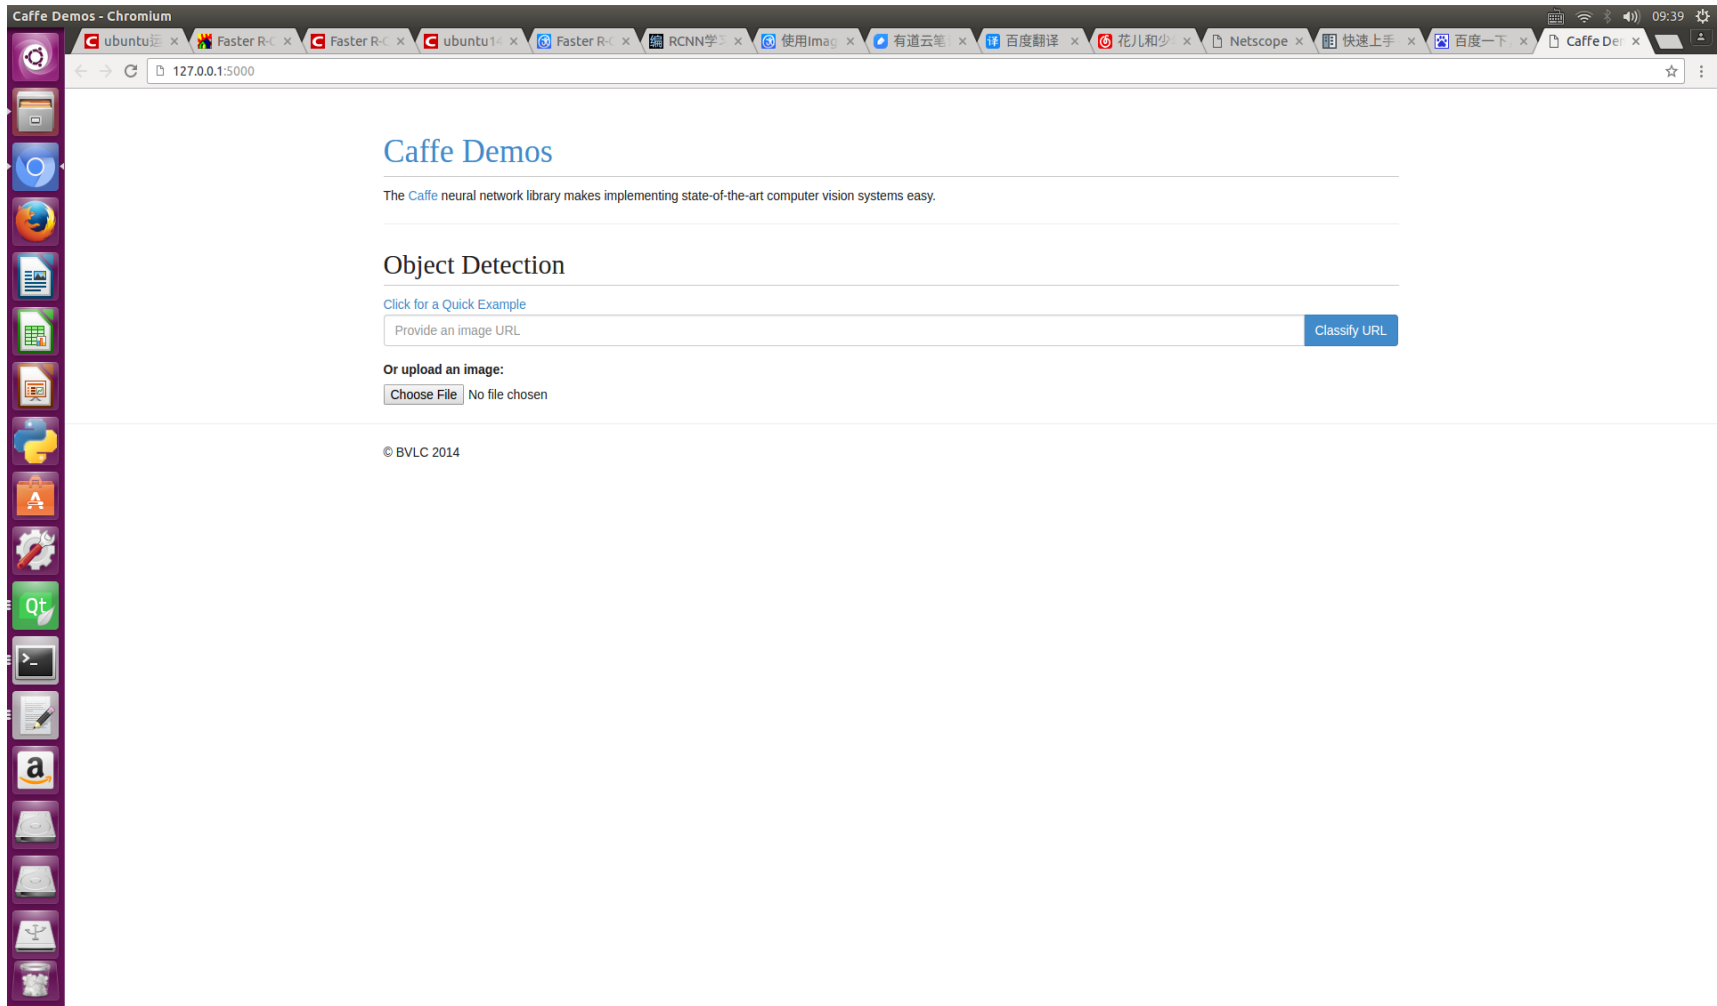
```
@app.route('/')
def index():
    return flask.render_template('index.html', has_result=False)
```

```
@app.route('/detection_url', methods=['GET'])
def dectection_url():
    imageurl = flask.request.args.get('imageurl', '')
    try:
        string_buffer = StringIO.StringIO(
            urllib.urlopen(imageurl).read())
        image = caffe.io.load_image(string_buffer)
```

```
except Exception as e:
```

Source: <http://docs.jinkan.org/docs/flask/quickstart.html#a-minimal-application>

Web Demos



Q & A